

Lecture 1

Introduction to Cryptography

Stefan Dziembowski

www.crypto.edu.pl/Dziembowski


University of Warsaw



Podstawowe informacje

- **Egzamin:** pisemny dwuczęściowy, kolokwium w trakcie semestru
- **Strona przedmiotu:** <http://www.crypto.edu.pl/teaching/crypto-i-1>
- **Podstawowy podręcznik:** Jonathan Katz and Yehuda Lindell
[Introduction to Modern Cryptography](#)
- **Pozostała literatura**
 - Doug Stinson [Cryptography Theory and Practice, Third Edition](#)
 - Shafi Goldwasser and Mihir Bellare [Lecture Notes on Cryptography](#)
 - Alfred J. Menezes, Paul C. van Oorschot and Scott A. Vanstone [Handbook of Applied Cryptography](#)

Plan




1. Introduction
2. Historical ciphers
3. Information-theoretic security
4. Computational security

Historical cryptography

cryptography \approx encryption

main applications: **military and diplomacy**



Modern cryptography

cryptography = much more
than encryption!



indistinguishability obfuscation

mental poker

signature schemes

electronic auctions

key agreement

e-cash

electronic voting

zero-knowledge

public-key
cryptography

multiparty-computations


sevenites

now

What happened in the seventies?


Technology

affordable
hardware



Demand

companies and
individuals start
to do business
electronically



Theory

**the
computational
complexity
theory is born**

this allows
researchers to
reason about
security in a
formal way.

Cryptography



In the past:

the **art** of encrypting messages (mostly for the military applications).

Now:

the **science** of securing digital communication and transactions (encryption, authentication, digital signatures, e-cash, auctions, etc..)



Three components of the course

1. practical aspects
2. mathematical foundations
3. new horizons

Practical aspects

- **symmetric encryption**: block ciphers (DES, AES) and stream ciphers (RC4)
- **hash functions** (MD5, SHA1,...), message authentication (CBC-MAC)
- **public-key infrastructure** (X.509, PGP, identity-based)
- **elements of number theory**
- **asymmetric encryption** (RSA, ElGamal, Rabin,...)
- **signature schemes** (RSA, ElGamal,...)

Mathematical foundations

- What makes us believe that the **protocols are secure**?
- Can we formally **define** “security”?
- Can security be **proven**?
- Do there exist “**unbreakable**” ciphers?

New horizons

Advanced cryptographic protocols, such as:

- **zero-knowledge**
- **multiparty computations**
- **private information retrieval**




This course is **not** about

- **practical data security** (firewalls, intrusion-detection, VPNs, etc.),

(however, we will talk a bit about the **cryptographic protocols** used in real life)

- **history** of cryptography,
- **number theory** and **algebra**
(we will use them **only as tools**)
- **complexity theory.**

Terminology



constructing secure
systems

breaking the systems

Cryptology = cryptography + cryptanalysis

This convention is **slightly artificial** and often ignored.

Common usage:

“cryptanalysis of X” = “breaking X”

Common abbreviation: **“crypto”**

Cryptography – general picture

plan of the course:

	encryption	authentication
private key	1 private key encryption	2 private key authentication
public key	3 public key encryption	4 signatures

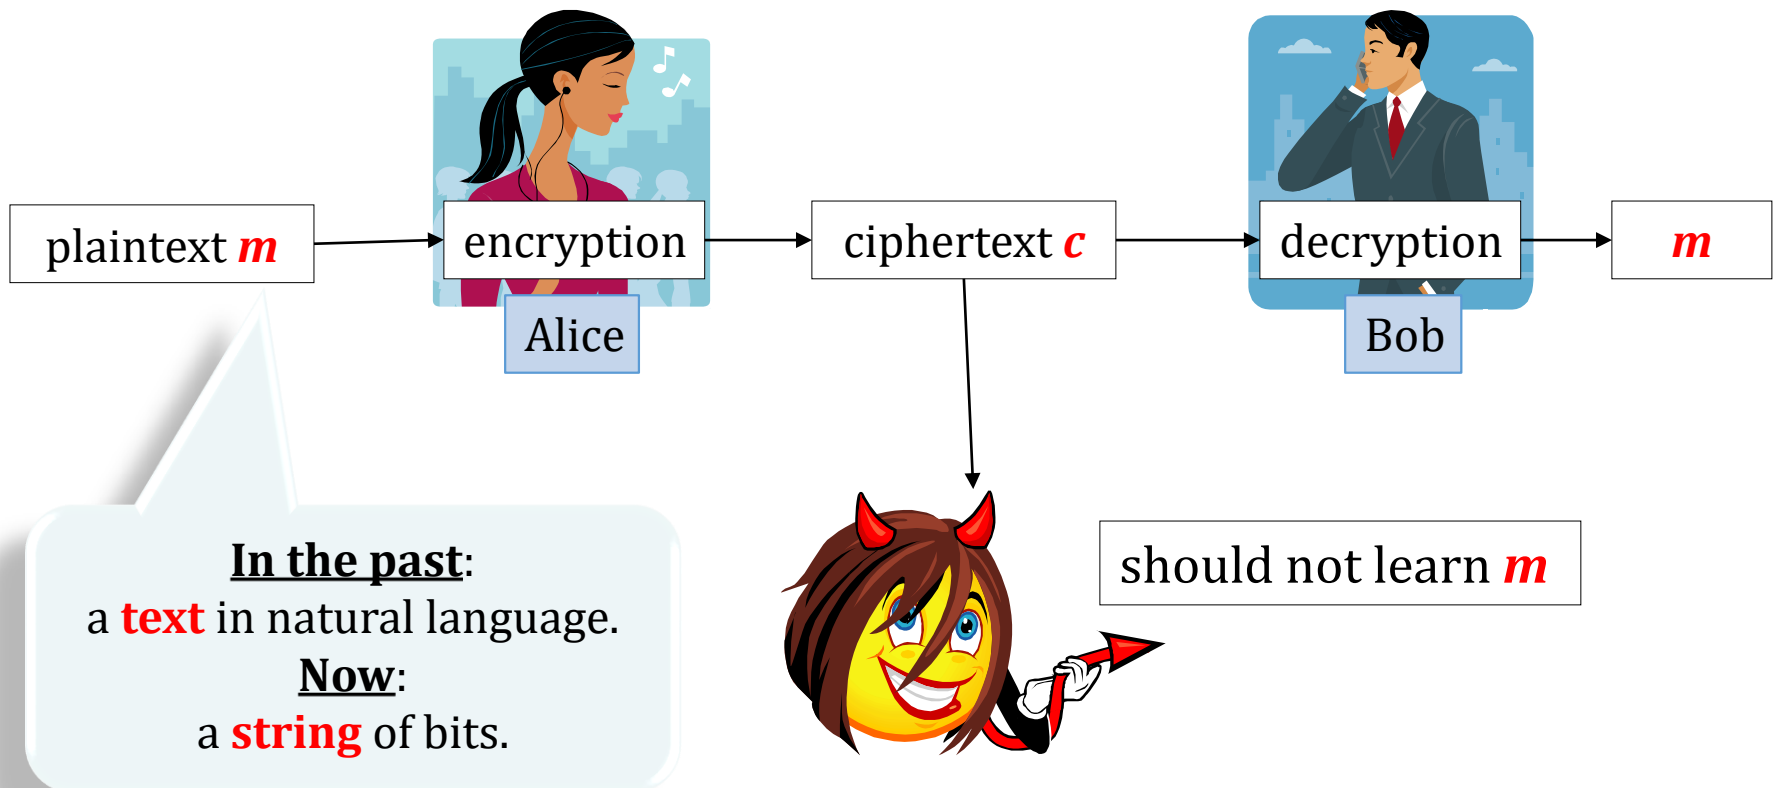
5 advanced cryptographic protocols

Preliminary plan of the lectures

1. Introduction to Cryptography
2. Symmetric Encryption I
3. Symmetric Encryption II
4. Symmetric Encryption III
5. Message Authentication and Introduction to Hash Functions
6. Key Management and Public-Key Cryptography
7. A Brush-up on Number Theory and Algebra
8. Public-Key Encryption I
9. Public-Key Encryption II
10. Signature Schemes and Commitment Schemes
11. Commitment Schemes and Zero-Knowledge Protocols
12. Two-party and Multi-party Computation Protocols
13. Private Information Retrieval
14. Introduction to Cryptographic Currencies

Encryption schemes (a very general picture)

Encryption scheme (cipher) = encryption & decryption



Art vs. science

In the past:

lack of precise definitions, ad-hoc design, usually insecure.

Nowadays:

formal definitions, systematic design, very secure constructions.

Provable security

We want to construct schemes that are
provably secure.

But...

- **why** do we want to do it?
- **how** to define it?
- and is it **possible** to achieve it?

Provable security – the motivation

In many areas of computer science formal proofs are **not essential**.

For example, instead of proving that an algorithm is efficient, we can just simulate it on a “*typical*” input”.

In **cryptography** it's **not true**, because

there cannot exist an experimental proof that a scheme is secure.

Why?

Because a notion of a

“*typical*” adversary”

does not make sense.

Security definitions are useful also because they allow us to construct schemes in a modular way...

Kerckhoffs' principle



Auguste Kerckhoffs (1883):

The enemy knows the system

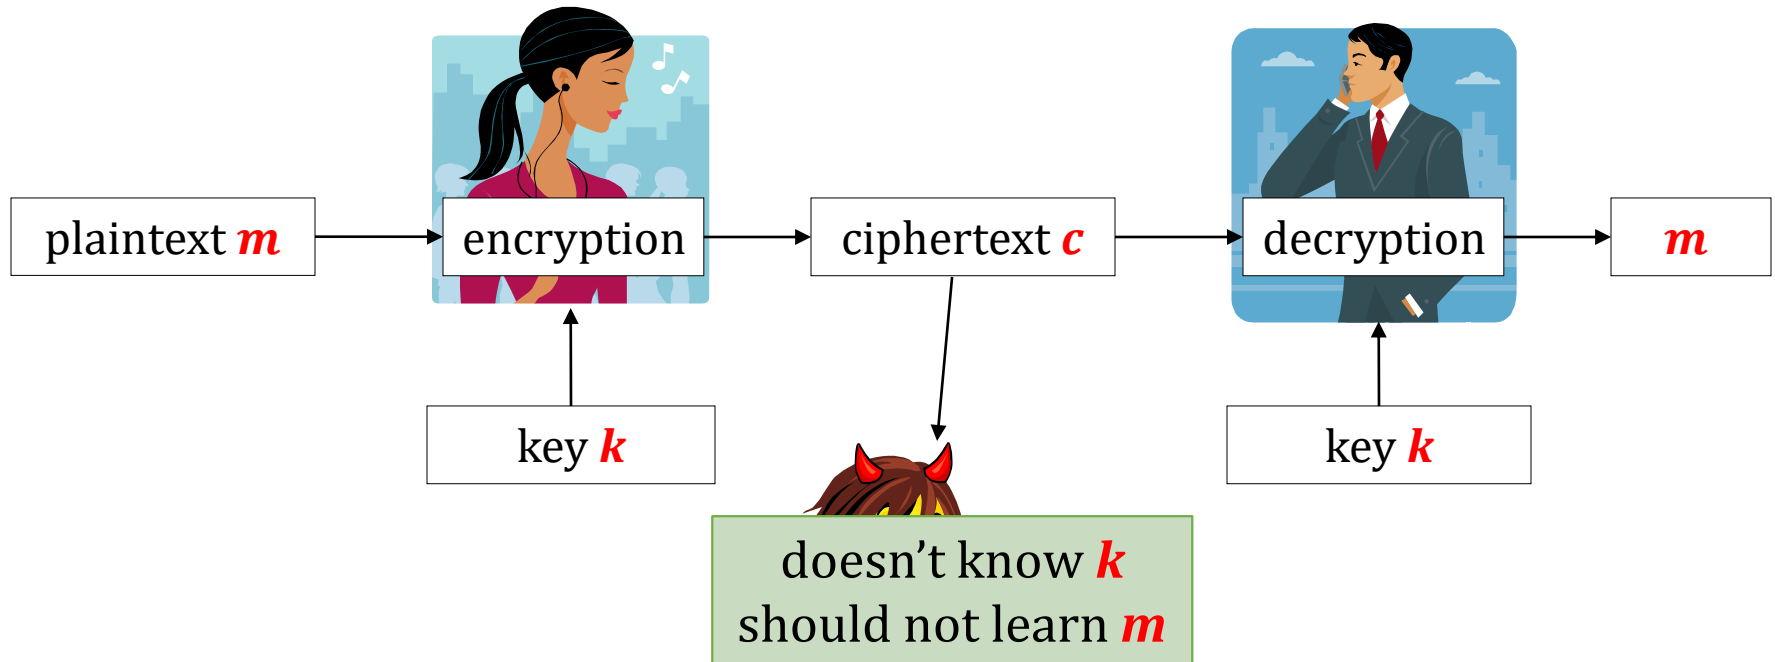
The cipher should remain secure even if **the adversary knows the specification of the cipher.**

The only thing that is **secret** is a

short key ***k***

that is **usually chosen uniformly at random**

A more refined picture



(Of course Bob can use the same method to send messages to Alice.)
(That's why it's called the **symmetric setting**)

Let us assume that k is uniformly random

Kerckhoffs' principle – the motivation

1. In commercial products it is unrealistic to assume that the design details remain secret (**reverse-engineering!**)
2. Short keys are easier to **protect, generate** and **replaced**.
3. The design details can be discussed and **analyzed in public**.

Not respecting this principle
=
`security by obscurity`.

A mathematical view

\mathcal{K} – key space

\mathcal{M} – plaintext space

\mathcal{C} – ciphertext space

An **encryption scheme** is a pair **(Enc, Dec)**, where

- **Enc** : $\mathcal{K} \times \mathcal{M} \rightarrow \mathcal{C}$ is an **encryption** algorithm,
- **Dec** : $\mathcal{K} \times \mathcal{C} \rightarrow \mathcal{M}$ is an **decryption** algorithm.


We will sometimes write **Enc_k(m)** and **Dec_k(c)** instead of **Enc(k,m)** and **Dec(k,c)**.

Correctness

for every **k** we should have **Dec_k(Enc_k(m)) = m**.

Plan

1. Introduction
2. Historical ciphers
3. Information-theoretic security
4. Computational security




Shift cipher

\mathcal{M} = words over alphabet $\{A, \dots, Z\} \approx \{0, \dots, 25\}$


$\mathcal{K} = \{0, \dots, 25\}$

$\text{Enc}_k(m_0, \dots, m_n) = (m_0 + k \bmod 26, \dots, m_n + k \bmod 26)$

$\text{Dec}_k(c_0, \dots, c_n) = (c_0 - k \bmod 26, \dots, c_n - k \bmod 26)$



Cesar: $k = 3$



Security of the shift cipher

How to break the shift cipher?

Check all possible keys!

Let c be a ciphertext.

For every $k \in \{0, \dots, 25\}$ check if $\text{Dec}_k(c)$ “makes sense”.

Most probably only one such k exists.

Thus $\text{Dec}_k(c)$ is the message.


This is called a **brute force attack**.

Moral: the key space needs to be large!

Substitution cipher

\mathcal{M} = words over alphabet $\{A, \dots, Z\} \approx \{0, \dots, 25\}$

\mathcal{K} = a set of permutations of $\{0, \dots, 25\}$



$$\text{Enc}_{\pi}(m_0, \dots, m_n) = (\pi(m_0), \dots, \pi(m_n))$$

$$\text{Dec}_{\pi}(c_0, \dots, c_n) = (\pi^{-1}(c_0), \dots, \pi^{-1}(c_n))$$

How to break the substitution cipher?

Use **statistical patterns** of the language.

For example: the **frequency tables**.

Texts of **50** characters can usually be broken this way.

Letter	Frequency
E	0.127
T	0.097
I	0.075
A	0.073
O	0.068
N	0.067
S	0.067
R	0.064
H	0.049
C	0.045
L	0.040
D	0.031
P	0.030
Y	0.027
U	0.024
M	0.024
F	0.021
B	0.017
G	0.016
W	0.013
V	0.008
K	0.008
X	0.005
Q	0.002
Z	0.001
J	0.001

Figure 7 - Frequency Table

Other famous historical ciphers

Vigenère cipher:



Blaise de Vigenère
(1523 - 1596)




Leon Battista Alberti
(1404 – 1472)

Enigma



Marian Rejewski
(1905 - 1980)



Alan Turing
(1912-1954)

In the past ciphers were designed in an ad-hoc manner


In contemporary cryptography the ciphers are designed in a **systematic way**.

Main goals:

1. define security
2. construct schemes that are “provably secure”

Plan

1. Introduction
2. Historical ciphers
3. Information-theoretic security
4. Computational security



Defining “security of an encryption scheme” is not trivial.

consider the following experiment

(m – a message)

1. the key K is chosen uniformly at random
2. $C := \text{Enc}_K(m)$ is given to the adversary

how to define
security



Idea 1

(m – a message)

1. the key K is chosen uniformly at random
2. $C := \text{Enc}_K(m)$ is given to the adversary

An idea

“The adversary should not be able to compute K .”

A problem

the encryption scheme that “doesn’t encrypt”:

$$\text{Enc}_K(m) = m$$

satisfies this definition!



Idea 2

(m – a message)

1. the key K is chosen uniformly at random
2. $C := \text{Enc}_K(m)$ is given to the adversary

An idea

“The adversary should not be able to compute m .”

A problem

What if the adversary can compute, e.g., the first half of m ?



Idea 3

(m – a message)

1. the key K is chosen uniformly at random
2. $C := \text{Enc}_K(m)$ is given to the adversary

An idea

“The adversary should not learn any information about m .”

A problem

But he may already have some a priori information about m !

For example he may know that m is a sentence in English...



Idea 4

(m – a message)

1. the key K is chosen uniformly at random
2. $C := \text{Enc}_K(m)$ is given to the adversary

An idea


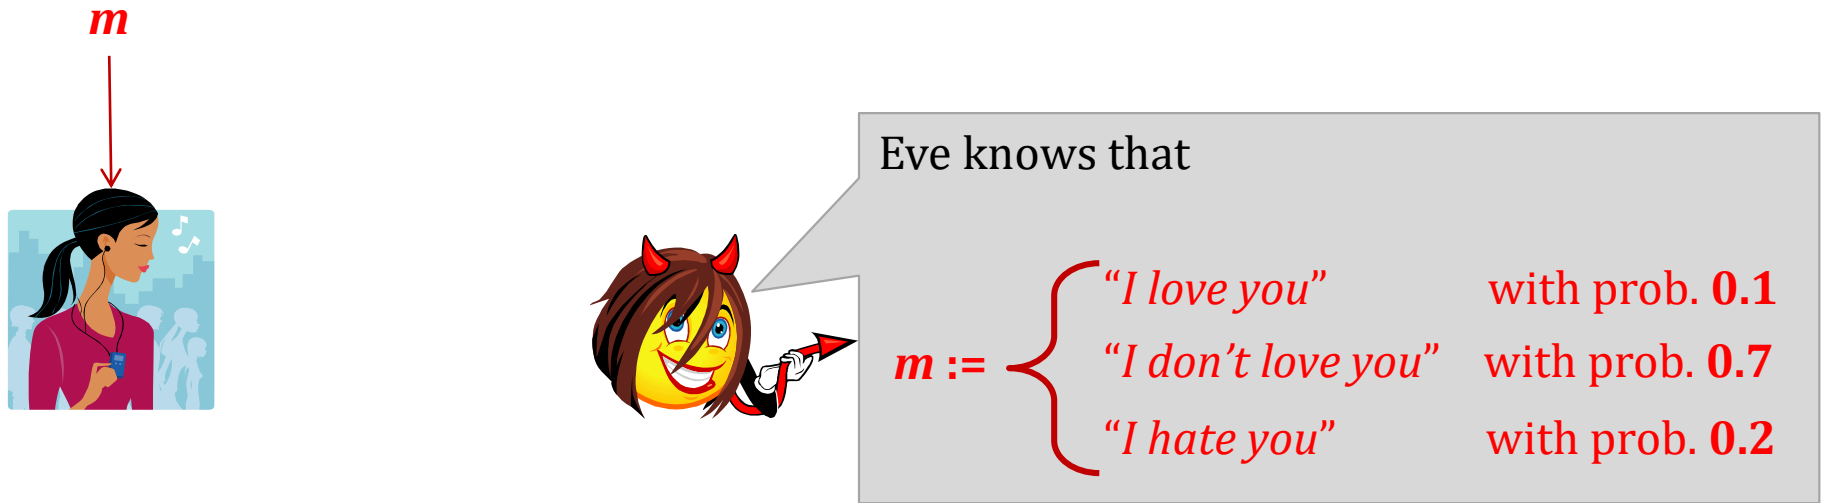
“The adversary should not learn any additional information about m .”

This makes much more sense.

But how to formalize it?



Example



How to formalize the “Idea 4”?

“The adversary should not learn any additional information about m .”

also called: **information-theoretically secret**

An encryption scheme is **perfectly secret** if

for every random variable M

and every $m \in \mathcal{M}$ and $c \in \mathcal{C}$

such that
 $P(C = c) > 0$

$$P(M = m) = P(M = m \mid (\text{Enc}(K, M)) = c)$$



equivalently: M and $\text{Enc}(K, M)$ are independent

Equivalently:

for every M we have that: M and $\text{Enc}(K,M)$ are independent



“the distribution of $\text{Enc}(K,m)$ does not depend on m ”



for every m_0 and m_1 we have that
 $\text{Enc}(K,m_0)$ and $\text{Enc}(K,m_1)$
have the same distribution

A perfectly secret scheme: one-time pad

t – a parameter
 $\mathcal{K} = \mathcal{M} = \{0,1\}^t$

component-wise **xor**

Vernam's cipher:

$$\text{Enc}_k(m) = k \text{ xor } m$$

$$\text{Dec}_k(c) = k \text{ xor } c$$



Gilbert
Vernam
(1890 –1960)

Correctness is trivial:

$$\text{Dec}_k(\text{Enc}_k(m)) = k \text{ xor } (k \text{ xor } m) \\ m$$

Perfect secrecy of the one-time pad

Perfect secrecy of the one time pad is also trivial.

This is because for every m
the distribution of $\text{Enc}(K, m)$ is uniform
(and hence does not depend on m).

for every c :

$$P(\text{Enc}(K, m) = c) = P(K = m \text{ xor } c) = 2^{-t}$$

Observation

One time pad can be **generalized** as follows.

Let $(G,+)$ be a group. Let $\mathcal{K} = \mathcal{M} = \mathcal{C} = G$.

The following is a perfectly secret encryption scheme:

- $\text{Enc}(k,m) = m + k$
- $\text{Dec}(k,m) = m - k$

Why the one-time pad is not practical?

1. The key has to be as long as the message.
2. The key cannot be reused

This is because:

$$\begin{aligned}\text{Enc}_k(m_0) \text{ xor } \text{Enc}_k(m_1) &= (k \text{ xor } m_0) \text{ xor } (k \text{ xor } m_1) \\ &= m_0 \text{ xor } m_1\end{aligned}$$



Theorem (Shannon 1949)

(“One time-pad is optimal in the class of perfectly secret schemes”)

In every perfectly secret encryption scheme

$$\text{Enc} : \mathcal{K} \times \mathcal{M} \rightarrow \mathcal{C}, \text{Dec} : \mathcal{K} \times \mathcal{C} \rightarrow \mathcal{M}$$

we have $|\mathcal{K}| \geq |\mathcal{M}|$.

Proof

Perfect secrecy implies that the distribution of $\text{Enc}(K, m)$ does not depend on m . Hence for every m_0 and m_1 we have

$$\{\text{Enc}(k, m_0)\}_{k \in \mathcal{K}} = \{\text{Enc}(k, m_1)\}_{k \in \mathcal{K}}$$

denote this set with \mathcal{C}'



Observation: $|\mathcal{K}| \geq |\mathcal{C}'|$.

Fact: we always have that $|\mathcal{C}'| \geq |\mathcal{M}|$.

This is because for every k we have that

$\text{Enc}_k : \mathcal{M} \rightarrow \mathcal{C}'$ is an injection

(otherwise we wouldn't be able to decrypt).




$$|\mathcal{K}| \geq |\mathcal{M}|$$

Practicality?

Generally, the **one-time pad** is **not very practical**, since:

- the key has to be as long as the **total** length of the encrypted messages,
- it is hard to generate truly random strings.



a **KGB** one-time pad hidden
in a walnut shell

However, it is sometimes used (e.g. in the **military applications**), because of the following advantages:

- **perfect secrecy**,
- short messages can be encrypted using **pencil and paper**.

In the 1960s the Americans and the Soviets established a hotline that was encrypted using the one-time pad. (**additional advantage**: they didn't need to share their secret encryption methods)

Venona project (1946 – 1980)



Ethel and Julius Rosenberg

American **National Security Agency** decrypted **Soviet** messages that were transmitted in the 1940s.

That was possible because the Soviets reused the keys in the one-time pad scheme.

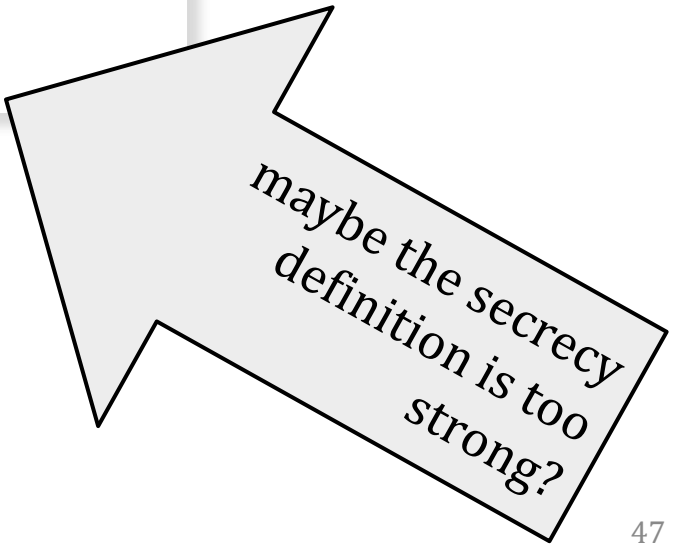
Outlook

We constructed a perfectly secret encryption scheme

Our scheme has certain drawbacks ($|\mathcal{K}| \geq |\mathcal{M}|$).

But by Shannon's theorem this is unavoidable.

Can we go home and relax?



maybe the secrecy
definition is too
strong?

What to do?

Idea

use a model where the **power** of the adversary is limited.

How?

Classical (computationally-secure) cryptography:

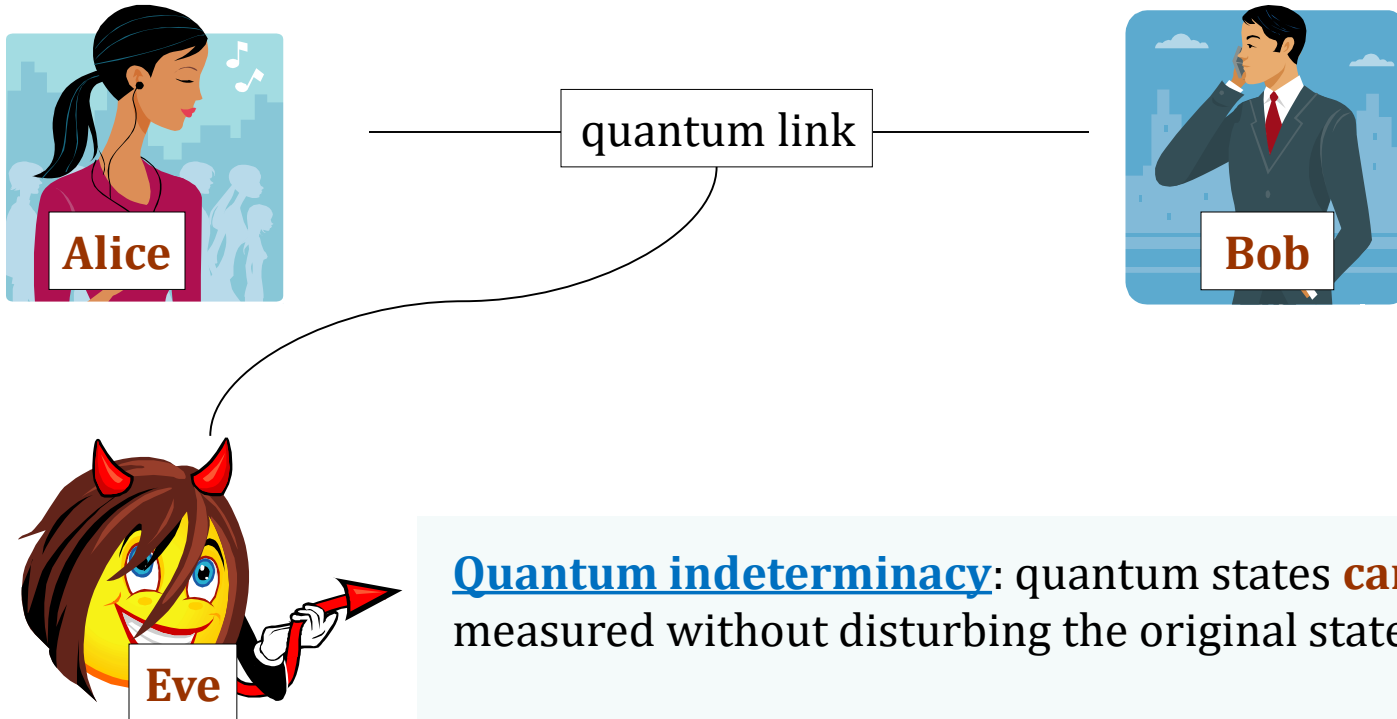
bound his computational power.

Alternative options:

quantum cryptography, bounded-storage model,...
(not too practical)

Quantum cryptography

Stephen Wiesner (1970s), Charles H. Bennett and Gilles Brassard (1984)



Quantum indeterminacy: quantum states **cannot** be measured without disturbing the original state.

Hence **Eve** cannot read the bits in an unnoticeable way.

Quantum cryptography

Advantage:

security is based on the laws of quantum physics

Disadvantage:

needs a dedicated equipment.

Practicality?

Currently: successful transmissions for distances of length around **150 km**.
Commercial products are available.


Warning:

Quantum **cryptography** should not be confused with quantum **computing**.

A satellite scenario

A third party (a satellite) is broadcasting random bits.

```
000110100111010010011010111001110111  
1110100111010101010010010010100111100  
001001111111100010101001000101010010  
0010100101001010111010101001010010101
```

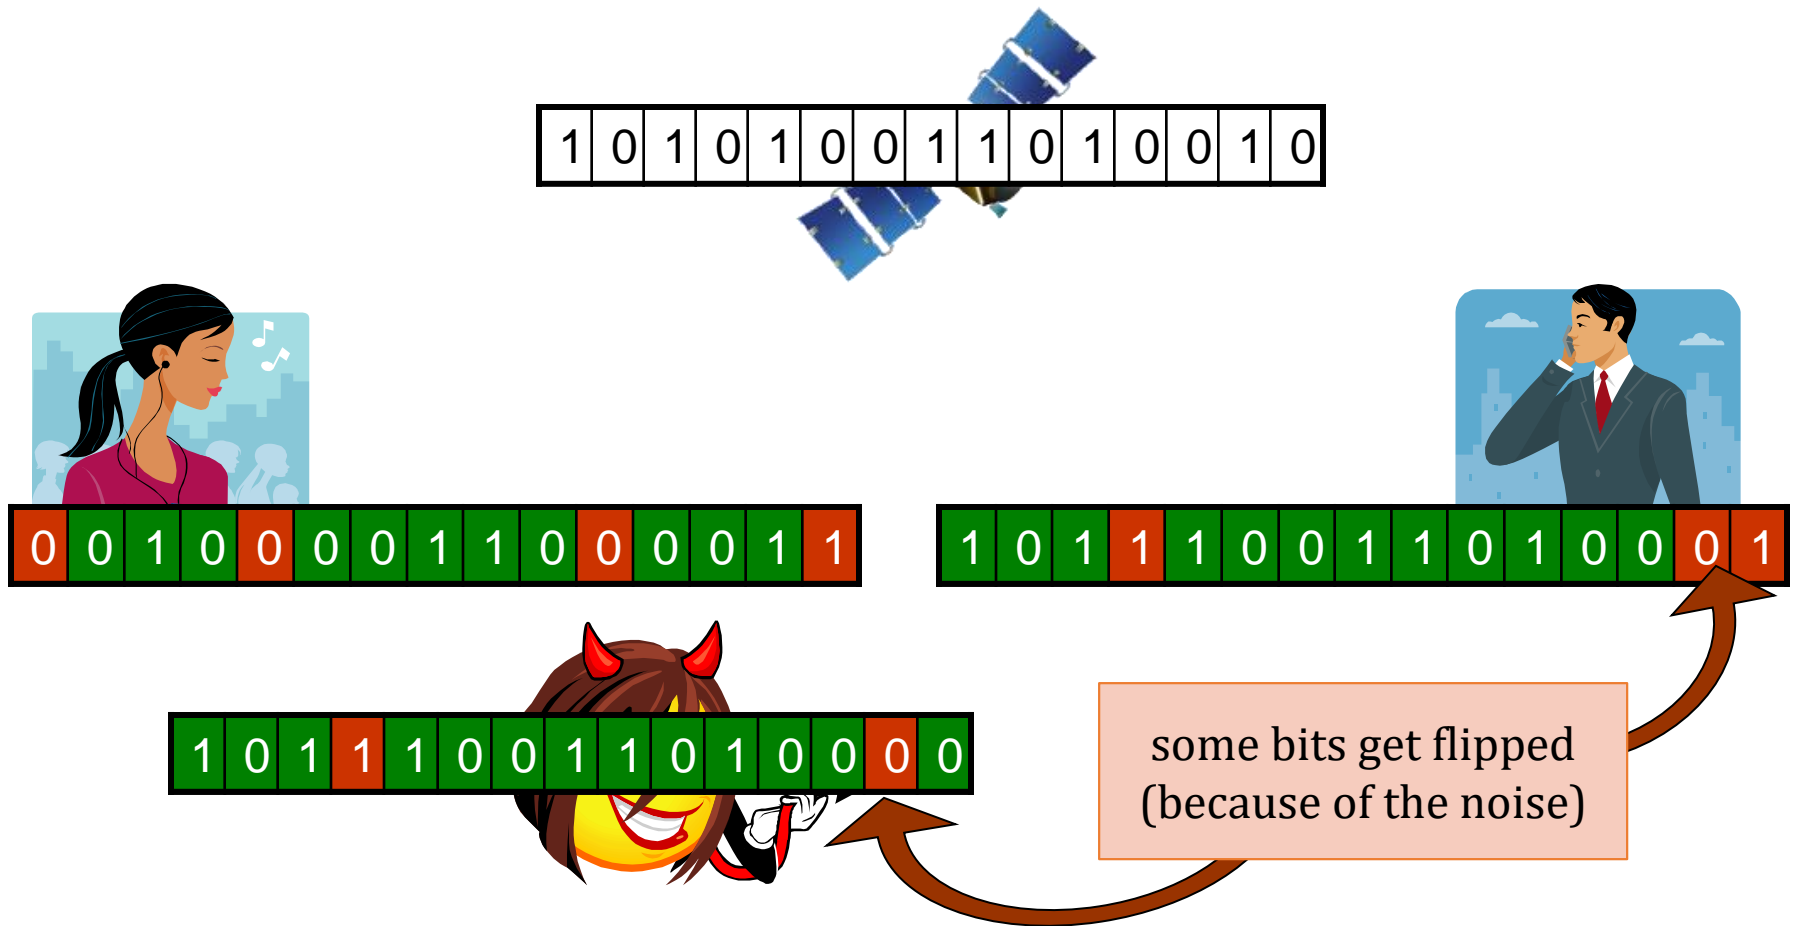


Does it help?

No...

(**Shannon's theorem** of course also holds in this case.)

Ueli Maurer (1993): noisy channel.



Assumption: the data that the adversary receives is noisy.
(The data that Alice and Bob receive may be even more noisy.)

Bounded-Storage Model

Another idea: bound the size of adversary's memory




```
000110100111010010011010111001110111
111010011101010101010010010100111100
001001111111100010101001000101010010
001010010100101011010101001010010101
```

← too large to fit in Eve's memory →



Plan

1. Introduction
2. Historical ciphers
3. Information-theoretic security
4. Computational security



How to reason about the bounded computing power?

perfect secrecy:
 M and $\text{Enc}_K(M)$
are independent

It is enough to require that

M and $\text{Enc}_K(M)$
are independent
“from the point of view of a computationally-limited adversary”.

How can this be formalized?

We will use the **complexity theory**!

Real cryptography starts here:



Eve is computationally-bounded

We will construct schemes that in **principle can be broken** if the adversary has a huge computing power.

For example, the adversary will be able to break the scheme by enumerating all possible secret keys.

(this is called a “**brute force attack**”)

Computationally-bounded adversary



Eve is computationally-bounded

But what does it mean?

Ideas:

1. “She has can use at most **1000**
Intel Core 2 Extreme X6800 Dual Core Processors
for at most **100** years...”
2. “She can buy equipment worth **1 million euro** and use it for **30** years..”.

it's hard to reason
formally about it

A better idea

"The adversary has access to a **Turing Machine** that can make at most **10^{30}** steps."

More generally, we could have definitions of a type:

"a system **X** is **(t, ϵ) -secure** if every **Turing Machine**
that operates in time **t**
can break it with probability at most **ϵ** ."

This would be quite precise, **but...**

We would need to specify exactly what we mean by a "**Turing Machine**":

- **how many tapes does it have?**
- how does it access these tapes (maybe a "**random access memory**" is a more realistic model..)
- ...

*Moreover, this approach often leads to **ugly formulas...***

What to do?

Idea

:

t steps of a Turing Machine \rightarrow “**efficient computation**”

$\epsilon \rightarrow$ a value “**very close to zero**”.

How to formalize it?

Use the **asymptotics**!

Efficiently computable?

“efficiently computable”

=

“polynomial-time computable
on a **Probabilistic Turing
Machine**”

that is: running in time
 $O(n^c)$ (for some **c**)

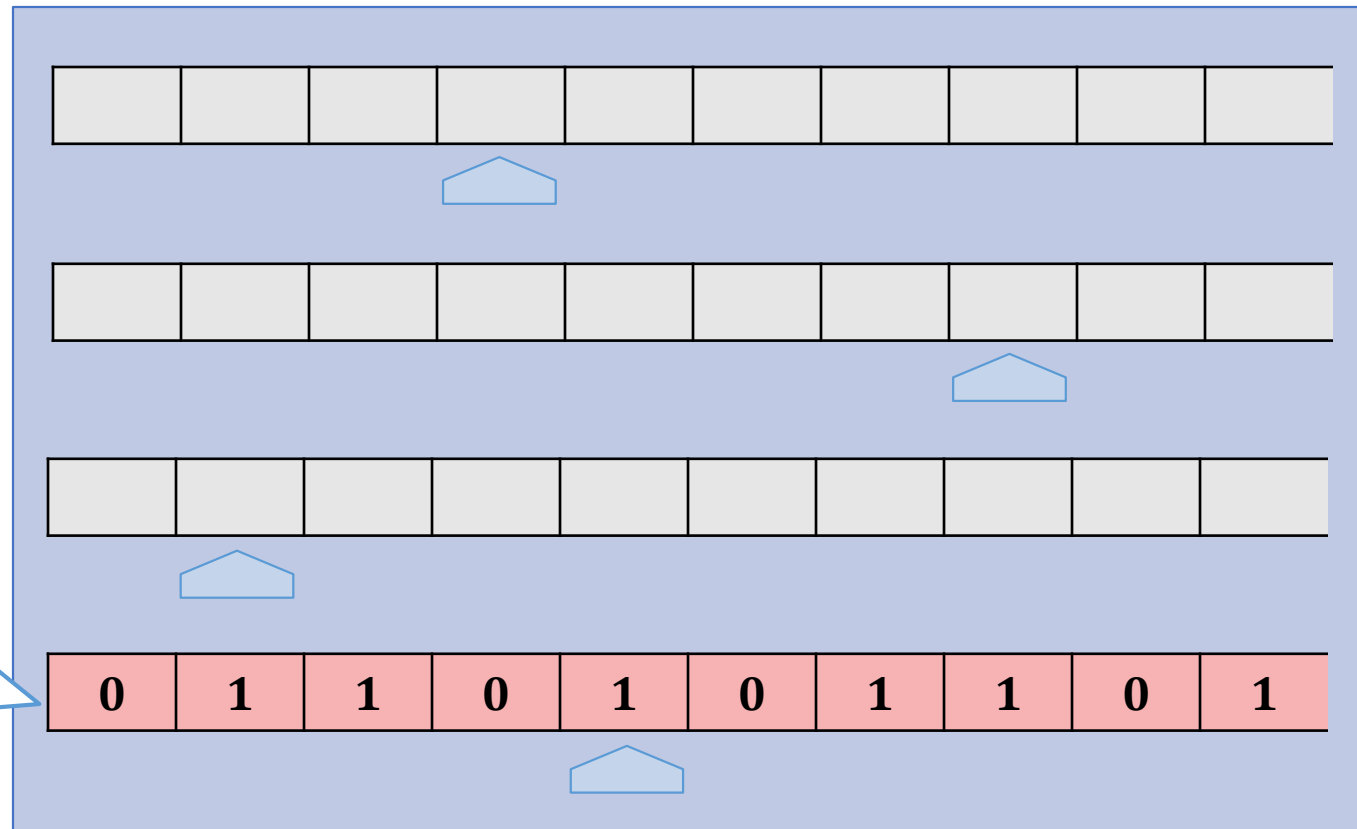
Here we assume that the **poly-time Turing Machines**
are the right model for the real-life computation.

Not true if a **quantum computer** is built...

Probabilistic Turing Machines

A standard Turing Machine has some number of tapes:

A **probabilistic** Turing Machine has an additional tape with random bits.



Some notation

If M is a Turing Machine then

$$M(X)$$

is a **random variable** denoting the **output** of M
assuming that
the contents of the random tape was chosen
uniformly at random.

More notation

$$Y \leftarrow M(X)$$

means that the variable Y takes the value that M outputs on input X (assuming the random input is chosen uniformly).

If \mathcal{A} is a set then

$$Y \leftarrow \mathcal{A}$$

means that Y is chosen uniformly at random from the set \mathcal{A} .

Very small?

“very small”

=

“negligible”

=

approaches **0** faster than the inverse of any polynomial

Formally

A function $\mu : \mathbf{N} \rightarrow \mathbf{R}$ is negligible if for every positive integer c there exists an integer N such that for all $x > N$

$$|\mu(x)| \leq \frac{1}{x^c}$$

Negligible or not?

$$f(n) := \frac{1}{n^2} \quad \text{no}$$

$$f(n) := 2^{-n} \quad \text{yes}$$

$$f(n) := 2^{-\sqrt{n}} \quad \text{yes}$$

$$f(n) := n^{-\log n} \quad \text{yes}$$

$$f(n) := \frac{1}{n^{1000}} \quad \text{no}$$

Nice properties of these notions

A sum of two polynomials is a polynomial:

$$\text{poly} + \text{poly} = \text{poly}$$

A product of two polynomials is a polynomial:

$$\text{poly} * \text{poly} = \text{poly}$$

A sum of two negligible functions is a negligible function:

$$\text{negl} + \text{negl} = \text{negl}$$

Moreover

A negligible function multiplied by a polynomial is negligible

$$\text{negl} * \text{poly} = \text{negl}$$

Security parameter

Typically, we will say that a scheme **X** is secure if



polynomial-time
Turing Machine **M**

P (**M** breaks the scheme **X**) is **negligible**

The terms “**negligible**” and “**polynomial**” make sense only if **X** and the **adversary** take an additional input **1ⁿ** called a **security parameter**.

In other words: we consider an infinite sequence **X(1),X(2),...** of schemes.

Example

security parameter n = the length of the secret key k

in other words: k is always a random element of $\{0,1\}^n$

The adversary can always **guess** k with probability 2^{-n} .

This probability **is negligible**.

He can also **enumerate all possible keys** k in time 2^n .
(the “brute force” attack)

This time **is exponential**.

Is this the right approach?

Advantages



1. All types of **Turing Machines** are “equivalent” up to a “**polynomial reduction**”.
Therefore we do need to specify the details of the model.
2. The formulas get much simpler.

Disadvantage



Asymptotic results don't tell us anything about security of the **concrete systems**.



However

Usually one can prove **formally** an asymptotic result and then argue **informally** that “the constants are reasonable”
(and can be calculated if one really wants).

How to change the security definition?

we will require that m_0, m_1 are chosen by a **poly-time adversary**

An encryption scheme is **perfectly secret** if for every $m_0, m_1 \in \mathcal{M}$
 $\text{Enc}(K, m_0)$ and $\text{Enc}(K, m_1)$ are identically distributed

we will require that no **poly-time adversary** can distinguish
 $\text{Enc}(K, m_0)$ from $\text{Enc}(K, m_1)$

A game

(Enc,Dec) – an encryption scheme



adversary
(polynomial-time probabilistic Turing machine)

security parameter
 1^n



oracle

chooses m_0, m_1 such that
 $|m_0| = |m_1|$

m_0, m_1

has to guess b

c

1. selects k randomly from $\{0,1\}^n$
2. chooses a random $b = 0,1$
3. calculates
 $c := \text{Enc}(k, m_b)$

Alternative name: **has indistinguishable encryptions**

Security definition:

We say that **(Enc,Dec)** is **semantically-secure** if any **polynomial time** adversary guesses b correctly with probability at most $\frac{1}{2} + \epsilon(n)$, where ϵ is negligible.

Testing the definition

Suppose the adversary can compute k from $\text{Enc}(k,m)$.
Can he win the game?

YES!

Suppose the adversary can compute **some bit of** m
from $\text{Enc}(k,m)$. Can he win the game?

YES!

Multiple messages

In real-life applications we need to encrypt **multiple messages with one key**.

The adversary may learn something about the key by looking at

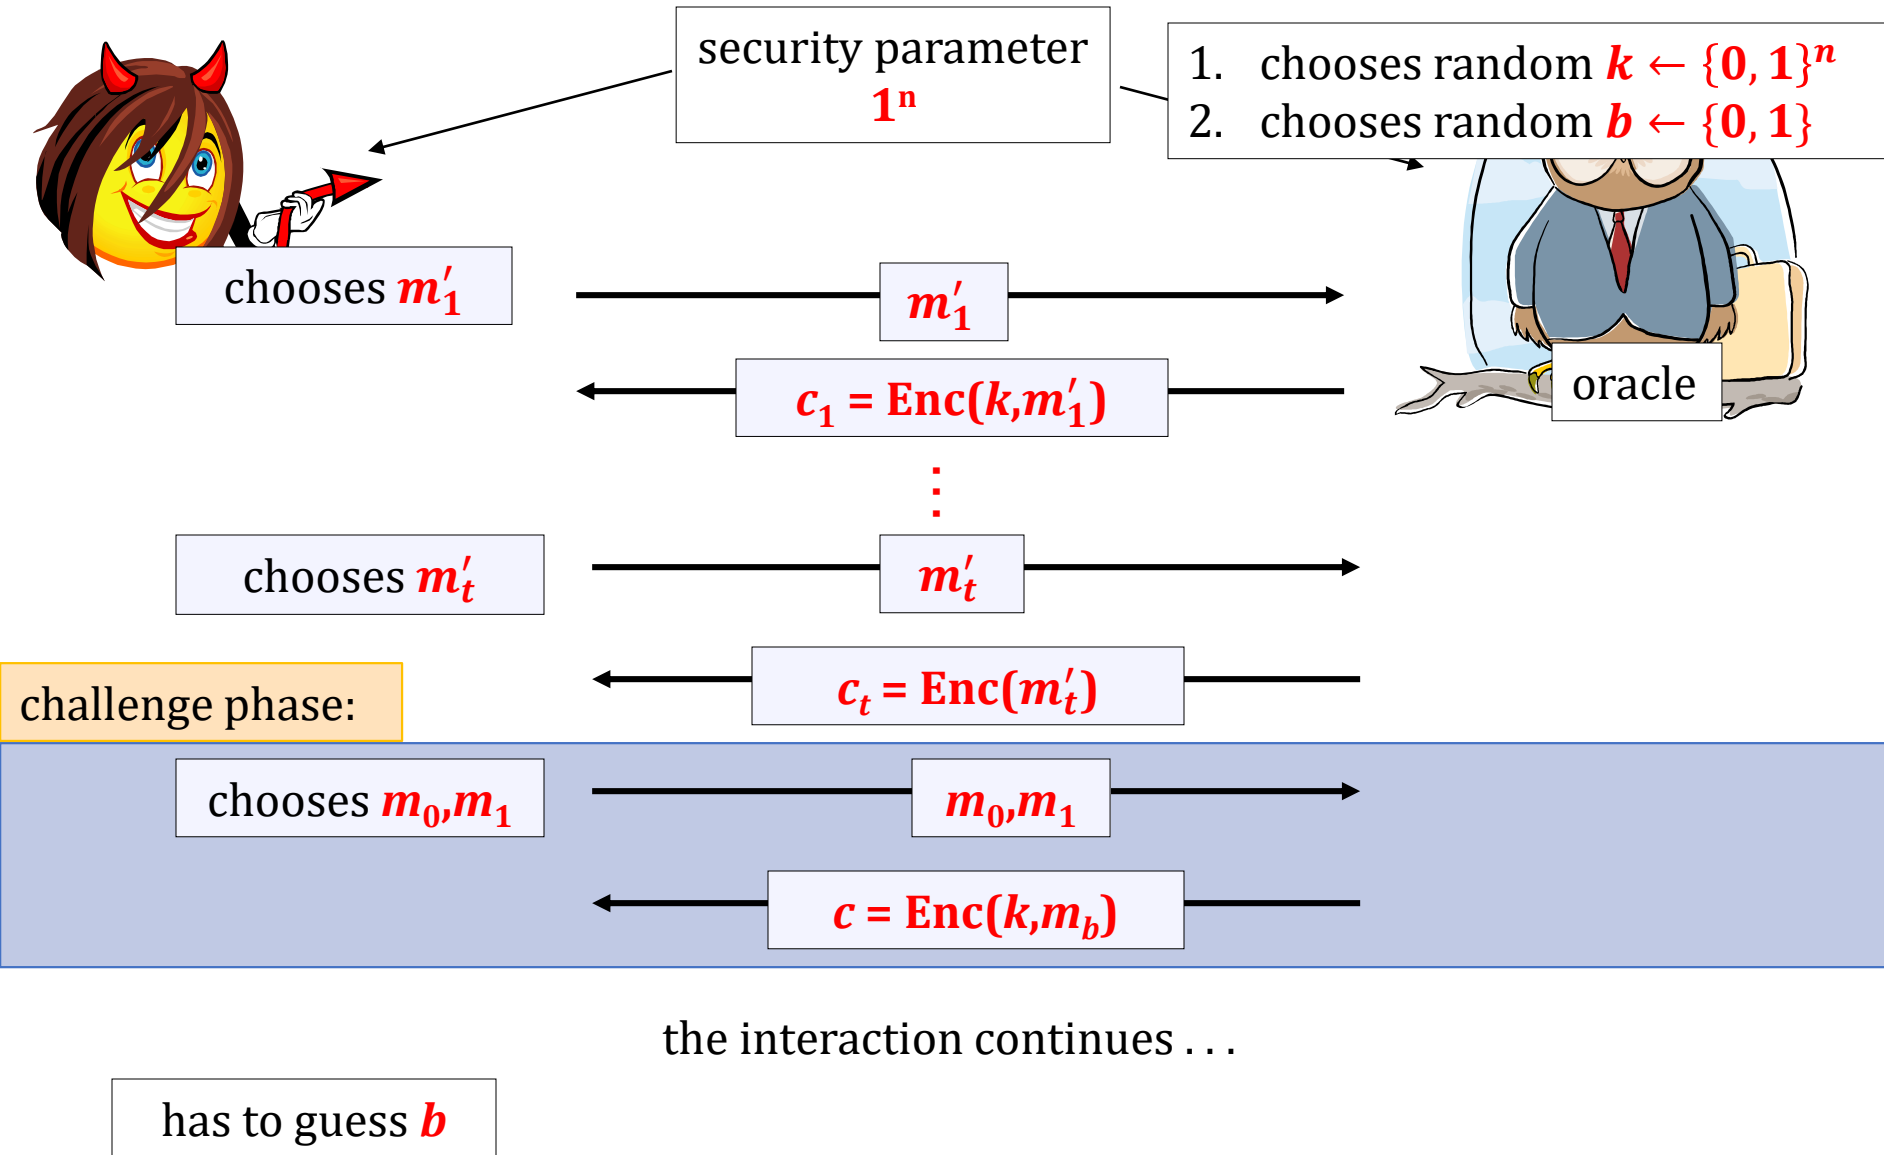
ciphertexts c_1, \dots, c_t of
some messages m_1, \dots, m_t .



**How are these messages chosen?
let's say: the adversary can choose them!**

(good tradition: be as pessimistic as possible)

A chosen-plaintext attack (CPA)



CPA-security

Alternative name: **CPA-secure**

Security definition

We say that **(Enc, Dec)** has indistinguishable encryptions under a chosen-plaintext attack (CPA) if

every **randomized polynomial time** adversary
guesses **b** correctly
with probability at most $\frac{1}{2} + \epsilon(n)$, where **ϵ** is negligible.

Observation

A **CPA-secure** encryption scheme cannot be deterministic.

Typical options:

- **Enc** has a “state” (e.g. a counter)
- **Enc** is randomized, i.e., it takes as additional input:
 - some perfect randomness ***R***, or
 - takes as an **nonce *R***

weaker
requirement


nonce = “**n**umber used **once**”

CPA in real-life

Q: Aren't we too pessimistic?

A: No! **CPA** can be implemented in practice.

Example: routing



Other attacks known in the literature



- **ciphertext-only attack** – the adversary has no information about the plaintext
- **known plaintext attack** – the plaintext are drawn from some distribution that the adversary does not control
- **batch chosen-plaintext attack** – like the **CPA** attack, but the adversary has to choose m_1, \dots, m_t at once.

(“our” **CPA**-attack is also called the “**adaptive CPA**-attack”)

- **chosen ciphertext attack** – we will discuss it later...

©2016 by Stefan Dziembowski. Permission to make digital or hard copies of part or all of this material is currently granted without fee *provided that copies are made only for personal or classroom use, are not distributed for profit or commercial advantage, and that new copies bear this notice and the full citation.*